

# CREDIT CARD FRAUD DETECTION

## Objective:

The primary objective of this case study is to build a machine learning model that can identify fraudulent credit card transactions. By analyzing the dataset, the model will distinguish between legitimate and fraudulent transactions. We will also explore various visualization techniques to understand the patterns of fraudulent activity, such as the distribution of transaction amounts and fraud occurrence by time

## Dataset Description:

- **Time:** Time of transaction in seconds since the first transaction.
- **V1-V28:** Anonymized features derived from PCA, representing transaction patterns.
- **Amount:** Monetary value of the transaction.
- **Class:** Target variable indicating transaction type:
  - 0: Legitimate transaction
  - 1: Fraudulent transaction

## Key Steps:

- **Load Data:** Import and explore dataset for shape, missing values, and basic stats.
- **Class Distribution:** Analyze and visualize the count of fraudulent vs. legitimate transactions.
- **Amount Statistics:** Calculate transaction amount statistics (min, max, mean, median).
- **Visualizations:** Plot transaction counts, correlations, and time distributions.
- **Transaction Hour Analysis:** Analyze and visualize fraudulent transactions by hour.
- **Hourly Stats:** Calculate transaction amounts by hour (sum, mean, etc.).
- **Amount Distribution:** Compare amounts for fraudulent vs. non-fraudulent transactions.
- **Scatter Plot:** Visualize fraudulent transactions' time vs. amount.

## Packages:

- **pandas:** Data manipulation and reading files.
- **numpy:** Numerical operations on arrays.
- **seaborn:** Statistical data visualization.
- **matplotlib.pyplot:** Plotting graphs and charts.
- **sklearn.model\_selection:** Splitting data into training and test sets.
- **sklearn.preprocessing:** Scaling data with StandardScaler.
- **sklearn.ensemble:** Random forest model for classification.

- **sklearn.metrics**: Evaluating models with classification reports and confusion matrix.
- **sklearn.linear\_model**: Logistic regression for binary classification.

## Step 1: Importing Libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.linear_model import LogisticRegression
```

## Step 2: Load the dataset

```
df = pd.read_csv('creditcard.csv')
print("Dataset loaded successfully:",df)

# Check for missing values and basic dataset info
print(f"Dataset shape: {df.shape}")
print("\nStatistical description of the dataset:\n",df.describe())
print(df.dtypes)
print(f"Missing values:\n{df.isnull().sum()}")
```

## Output:

Dataset loaded successfully:

Dataset loaded successfully:			Time	V1	V2	V3	V4	V5	V6	...	V24	V25	V26	V27	
V28	Amount	Class													
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	...	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	...	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	...	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	...	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	...	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	...	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	0
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	...	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	...	0.640134	0.265745	-0.087371	0.004455	-0.026561	67.88	0
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	...	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	0
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	...	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	0

Dataset shape: (284807, 31)

Statistical description of the dataset:

Statistical description of the dataset:											
	Time	V1	V2	V3	V4	...	V26	V27	V28	Amount	Class
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	...	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000	284807.000000
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	...	1.683437e-15	-3.660091e-16	-1.227390e-16	88.349619	0.001727
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	...	4.822270e-01	4.036325e-01	3.300833e-01	250.120109	0.041527
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	...	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000	0.000000
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	...	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000	0.000000
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	...	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000	0.000000
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	...	2.409522e-01	9.104512e-02	7.827995e-02	77.165000	0.000000
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	...	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000	1.000000

[8 rows x 31 columns]

**datatypes:** Time    float64

V1    float64  
 V2    float64  
 V3    float64  
 V4    float64  
 V5    float64  
 V6    float64  
 V7    float64  
 V8    float64  
 V9    float64  
 V10   float64  
 V11   float64  
 V12   float64  
 V13   float64  
 V14   float64  
 V15   float64  
 V16   float64  
 V17   float64  
 V18   float64  
 V19   float64  
 V20   float64  
 V21   float64  
 V22   float64  
 V23   float64  
 V24   float64  
 V25   float64  
 V26   float64  
 V27   float64  
 V28   float64  
 Amount   float64  
 Class    int64

#### Missing values:

Time	0
V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
Amount	0
Class	0

### Step 3: Analyze transaction types (fraud vs. non-fraud)

```
# Count the number of fraudulent and legitimate transactions
transaction_counts = df['Class'].value_counts()
print("Number of legitimate transactions:", transaction_counts[0])
print("Number of fraudulent transactions:", transaction_counts[1])

# Calculate the percentage of fraudulent transactions
total_transactions = len(df)
transaction_counts = df['Class'].value_counts()
fraudulent_transactions = transaction_counts[1]
```

```
non_fraudulent_transactions = transaction_counts[0]
percentage_fraudulent = (fraudulent_transactions / total_transactions) * 100
print(f"Percentage of fraudulent transactions: {percentage_fraudulent:.2f}%")
```

**Output:**


**Number of legitimate transactions: 284315**

**Number of fraudulent transactions: 492**

**Percentage of fraudulent transactions: 0.17%**

**Step 4 : Visualize distribution of fraudulent vs non-fraudulent transactions**

```
Data for the pie chart
labels = ['Non-Fraudulent', 'Fraudulent']
sizes = [non_fraudulent_transactions, fraudulent_transactions]
plt.pie(sizes, labels=labels, autopct='%2.2f%%', startangle=140,
colors=['#ff9999', '#66b3ff'])
plt.title('Distribution of Fraudulent Transactions')
plt.show()
```



## Step 5: Basic statistics on the 'Amount' column

```
# Calculate descriptive statistics for the 'Amount' column
min_amount = df['Amount'].min()
max_amount = df['Amount'].max()
mean_amount = df['Amount'].mean()
median_amount = df['Amount'].median()

# Print the statistics
print(f"Minimum amount: {min_amount:.2f}")
print(f"Maximum amount: {max_amount:.2f}")
print(f"Mean amount: {mean_amount:.2f}")
print(f"Median amount: {median_amount:.2f}")
```

### Output

Minimum amount: 0.00

Maximum amount: 25691.16

Mean amount: 88.35

Median amount: 22.00

## Step 6: Analyze the maximum transaction amount

```
# Find the maximum transaction amount
max_amount = df['Amount'].max()
max_amount_row = df[df['Amount'] == max_amount]
is_fraudulent = max_amount_row['Class'].values[0]
print(f"The maximum transaction amount is {max_amount}, and it is {'fraudulent' if is_fraudulent else 'legitimate'}.")
```


### Output:

The maximum transaction amount is 25691.16, and it is legitimate.

## Step 7: Count plot for fraudulent vs legitimate transactions


```
#Bar Chart Showing the Count of Fraudulent vs. Legitimate Transactions
plt.figure(figsize=(6, 4))
sns.countplot(x='Class', data=df)
plt.title('Count of Fraudulent vs. Legitimate Transactions')
```

```
plt.xlabel('Class (0: Legitimate, 1: Fraudulent)')
plt.ylabel('Count')
plt.show()
```



## Step 8: Correlation matrix heatmap

```
#Heatmap to Visualize the Correlation Between Numerical Features
correlation_matrix=df.corr()
plt.figure(figsize=(10,6))
sns.heatmap(correlation_matrix,cmap='coolwarm',vmin=-1,vmax=1,annot=False,linewidths=0.5)
plt.title('Correlation Matrix Heatmap')
plt.show()
```



## Step 9: Time-related analysis: Plot fraudulent vs non-fraudulent transaction times


```
# Ensure there are no infinite values in the 'Time' column
df['Time'].replace([float('inf'), -float('inf')], float('nan'), inplace=True)

# Splitting the data
class_0 = df.loc[df['Class'] == 0]["Time"]
class_1 = df.loc[df['Class'] == 1]["Time"]

# Creating a Seaborn KDE plot
plt.figure(figsize=(10, 6))
sns.kdeplot(class_0, label='Not Fraud', fill=True, common_norm=False)
sns.kdeplot(class_1, label='Fraud', fill=True, common_norm=False)

# Adding titles and labels
plt.title('Credit Card Transactions Time Density Plot')
plt.xlabel('Time [s]')
plt.ylabel('Density')
plt.legend()
plt.show()
```





## Step 10: Hourly analysis of fraudulent transactions

```
#Hour with the Most Frequent Fraudulent Transactions
# Convert the 'Time' column to hours
df['transaction_hour'] = (df['Time'] // 3600) % 24

# Assuming 'Class' column indicates fraudulent transactions (1 for fraud, 0 for non-fraud)
df['is_fraudulent'] = df['Class']

# Separate the fraudulent and non-fraudulent transactions
fraudulent_df = df[df['is_fraudulent'] == 1]
non_fraudulent_df = df[df['is_fraudulent'] == 0]
# Aggregate by hour for fraudulent and non-fraudulent transactions
fraudulent_by_hour = fraudulent_df.groupby('transaction_hour').size()
non_fraudulent_by_hour = non_fraudulent_df.groupby('transaction_hour').size()

# Plotting
plt.figure(figsize=(10, 6))
plt.bar(fraudulent_by_hour.index - 0.2, fraudulent_by_hour.values, width=0.4,
color='red', alpha=0.7, label='Fraudulent')
plt.bar(non_fraudulent_by_hour.index + 0.2, non_fraudulent_by_hour.values,
width=0.4, color='blue', alpha=0.7, label='Non-Fraudulent')
plt.xlabel('Hour of Day')
```

```

plt.ylabel('Number of Transactions')
plt.title('Transactions by Hour of Day')
plt.xticks(range(0, 24)) # Set x-ticks for hours
plt.legend()
plt.tight_layout()
plt.show()

# Convert the 'Time' column to hours
df['transaction_hour'] = (df['Time'] // 3600) % 24


# Assuming 'Class' column indicates fraudulent transactions (1 for fraud, 0 for non-fraud)
df['is_fraudulent'] = df['Class']


# Separate the fraudulent transactions
fraudulent_df = df[df['is_fraudulent'] == 1]

# Aggregate by hour for fraudulent transactions
fraudulent_by_hour = fraudulent_df.groupby('transaction_hour').size()
# Find the hour with the most frequent fraudulent transactions
most_frequent_hour = fraudulent_by_hour.idxmax()
most_frequent_count = fraudulent_by_hour.max()

# Print the result
print(f"The hour with the most frequent fraudulent transactions is: {most_frequent_hour}:00")
print(f"Number of fraudulent transactions during this hour: {most_frequent_count}")

```





The hour with the most frequent fraudulent transactions is: 2.0:00

Number of fraudulent transactions during this hour: 57

**Step 11: 10. Compare transaction statistics (Total, Mean, Max, Median) by hour for fraudulent and non-fraudulent transactions**

```
# Convert the 'Time' column to hours
df['transaction_hour'] = (df['Time'] // 3600) % 24

# Replace inf values with NaN
df.replace([np.inf, -np.inf], np.nan, inplace=True)

# Calculate required statistics for each hour
hourly_stats = df.groupby('transaction_hour').agg({
    'Amount': ['sum', 'mean', 'min', 'max', 'median']
}).reset_index()

# Rename columns for easier access
hourly_stats.columns = ['Hour', 'Total', 'Mean', 'Min', 'Max', 'Median']
# Separate fraudulent and non-fraudulent transactions
fraudulent_df = df[df['Class'] == 1]
non_fraudulent_df = df[df['Class'] == 0]

# Calculate required statistics for fraudulent transactions
fraudulent_stats = fraudulent_df.groupby('transaction_hour').agg({
    'Amount': ['sum', 'mean', 'min', 'max', 'median']
}).reset_index()
```

```

fraudulent_stats.columns = ['Hour', 'Total', 'Mean', 'Min', 'Max', 'Median']

# Calculate required statistics for non-fraudulent transactions
non_fraudulent_stats = non_fraudulent_df.groupby('transaction_hour').agg({
    'Amount': ['sum', 'mean', 'min', 'max', 'median']
}).reset_index()
non_fraudulent_stats.columns = ['Hour', 'Total', 'Mean', 'Min', 'Max', 'Median']

```

## Step 12: Plotting total and maximum transaction amounts

```


# Set up the plotting environment
sns.set(style="whitegrid")

# Plot Total Amount
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(18, 6))
sns.lineplot(ax=ax1, x='Hour', y='Total', data=non_fraudulent_stats, label='Non-
Fraudulent')
sns.lineplot(ax=ax2, x='Hour', y='Total', data=fraudulent_stats,
label='Fraudulent', color='red')
ax1.set_title('Total Amount of Non-Fraudulent Transactions by Hour')
ax2.set_title('Total Amount of Fraudulent Transactions by Hour')
plt.suptitle('Total Amount of Transactions by Hour')
plt.show()


# Plot Maximum Amount
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(18, 6))
sns.lineplot(ax=ax1, x='Hour', y='Max', data=non_fraudulent_stats, label='Non-
Fraudulent')
sns.lineplot(ax=ax2, x='Hour', y='Max', data=fraudulent_stats,
label='Fraudulent', color='red')
ax1.set_title('Maximum Amount of Non-Fraudulent Transactions by Hour')
ax2.set_title('Maximum Amount of Fraudulent Transactions by Hour')
plt.suptitle('Maximum Amount of Transactions by Hour')
plt.show()

```

Total Amount of Transactions by Hour



Maximum Amount of Transactions by Hour



```
# Create a temporary DataFrame with 'Amount' and 'Class' columns
tmp = df[['Amount', 'Class']].copy()

# Separate the data into non-fraudulent and fraudulent transactions
class_0 = tmp.loc[tmp['Class'] == 0]['Amount']
class_1 = tmp.loc[tmp['Class'] == 1]['Amount']

# Display summary statistics
print("Non-Fraudulent Transactions:", class_0.describe())
print("\nFraudulent Transactions:", class_1.describe())
```

```

# Filter data for fraudulent transactions
fraud = df.loc[df['Class'] == 1]
print("data for fraudulent transactions:", fraud)

# Create the scatter plot using Seaborn
plt.figure(figsize=(10, 6))
sns.scatterplot(x=fraud['Time'], y=fraud['Amount'], color='red', alpha=0.5)
plt.title('Amount of Fraudulent Transactions')
plt.xlabel('Time [s]')
plt.ylabel('Amount')
plt.show()

```

### Output:

**Non-Fraudulent Transactions: count 284315.000000**

mean 88.291022

std 250.105092

min 0.000000

25% 5.650000

50% 22.000000

75% 77.050000

max 25691.160000

Name: Amount, dtype: float64

**Fraudulent Transactions: count 492.000000**

mean 122.211321

std 256.683288

min 0.000000

25% 1.000000


50% 9.250000

75% 105.890000

max 2125.870000

#### data for fraudulent transactions:

data for fraudulent transactions:													Time	V1	V2	V3	V4	V5	...	V27	V28	Amount	Class	transaction
_hour	is_fraudulent																							
541	406.0	-2.312227	1.951992	-1.609851	3.997906	-0.522188	...	0.261145	-0.143276	0.00	1										0.0	1		
623	472.0	-3.043541	-3.157307	1.088463	2.288644	1.359805	...	-0.252773	0.035764	529.00	1										0.0	1		
4920	4462.0	-2.303350	1.759247	-0.359745	2.330243	-0.821628	...	0.039566	-0.153029	239.93	1										1.0	1		
6108	6986.0	-4.397974	1.358367	-2.592844	2.679787	-1.128131	...	-0.827136	0.849573	59.00	1										1.0	1		
6329	7519.0	1.234235	3.019740	-4.304597	4.732795	3.624201	...	-0.010016	0.146793	1.00	1										2.0	1		
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	...	0.292680	0.147968	390.00	1										22.0	1		
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	...	0.389152	0.186637	0.76	1										23.0	1		
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	...	0.385107	0.194361	77.89	1										23.0	1		
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	...	0.884876	-0.253700	245.00	1										23.0	1		
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	...	0.002988	-0.015309	42.53	1										23.0	1		



#### Conclusion:

The analysis of credit card transactions highlighted the rarity of fraudulent transactions, which account for only 0.17% of the total dataset. Despite their low occurrence, fraudulent transactions exhibit distinct patterns, such as a higher frequency during late hours (12 AM to 3 AM) and often involving higher transaction amounts. Visualizations revealed that these transactions tend to be more concentrated in specific hours, and high-value fraudulent transactions occur within shorter time intervals. To improve fraud detection, it is crucial to address the class imbalance through techniques like oversampling or under sampling. Additionally, developing a machine learning model that leverages time-based features and transaction amounts could enhance detection accuracy. Evaluating model performance with metrics like precision, recall, and F1-score will be key to ensuring its effectiveness in identifying fraud.

